

Practical Assignment 8: Decision Trees on Diabetes Dataset

Submission Details

Field	Details
Name	Ayushkar Pau
ID	GF202343142
Subject	Statistical Foundation of Data Science (CSU1658)
Date	December 9, 2025
Repo	View My Repository

Objective

The goal of this assignment is to build a **Decision Tree Classifier** to predict the onset of diabetes based on diagnostic measures. We will:

1. **Prepare the Data:** Load the **Pima Indians Diabetes Dataset** and perform feature selection (splitting target and independent variables).
 2. **Build the Model:** Train a Decision Tree using Scikit-Learn.
 3. **Visualize:** Create a pictorial representation of the decision tree structure.
 4. **Analyze Splits:** Calculate **Entropy**, **Information Gain**, and **Gini Index** to mathematically justify the selection of the root node.
-

Environment Setup and Dependencies

Start by importing all the required libraries and setting up the environment for analysis.

```
# --- 0. Environment Setup ---
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn import metrics

# Configure plotting
plt.rcParams['figure.figsize'] = (12, 8)
```

```

sns.set_style("whitegrid")
print("Environment setup complete.")
Environment setup complete.

```

Step 1: Data Loading and Feature Selection

Instruction: Divide given columns into two types of variables: dependent (target) and independent (features).

Approach

We will load the dataset and separate the columns.

- **Target Variable (y):** Outcome (0 = No Diabetes, 1 = Diabetes).
- **Feature Variables (X):** All other columns (Glucose, BMI, Age, etc.).

```

# --- Step 1: Load Data & Feature Selection ---

# Load Pima Indians Diabetes dataset
# (Using a direct raw URL so this runs anywhere without downloading a
# file manually)
url =
"https://raw.githubusercontent.com/jbrownlee/Datasets/master/pima-
indians-diabetes.data.csv"
col_names = ['Pregnancies', 'Glucose', 'BloodPressure',
'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age',
'Outcome']
df = pd.read_csv(url, header=None, names=col_names)

# Display first few rows
print("--- Dataset Head ---")
display(df.head())

# Feature Selection
feature_cols = ['Pregnancies', 'Glucose', 'BloodPressure',
'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']
X = df[feature_cols] # Features
y = df['Outcome']     # Target variable

print(f"\nFeature Matrix Shape: {X.shape}")
print(f"Target Vector Shape: {y.shape}")

--- Dataset Head ---

    Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin
BMI \
0            6        148             72              35      0   33.6
1            1         85              66              29      0   26.6

```

2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Feature Matrix Shape: (768, 8)
 Target Vector Shape: (768,)

Step 2: Split Data and Build Decision Tree

Instruction: Split the data and then build decision tree.

Approach

- **Split:** We use `train_test_split` (80% training, 20% testing) to ensure we can evaluate the model on unseen data.
- **Model:** We initialize the `DecisionTreeClassifier`.
 - We will use `criterion='entropy'` explicitly to help with our later calculations of Information Gain.
 - We set `max_depth=3` to keep the tree visualization readable and prevent overfitting.

```
# --- Step 2: Train-Test Split & Model Building ---

# 1. Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# 2. Create Decision Tree classifier object
# Using 'entropy' to optimize for Information Gain
clf = DecisionTreeClassifier(criterion="entropy", max_depth=3,
random_state=42)

# 3. Train Decision Tree Classifier
clf = clf.fit(X_train, y_train)

# 4. Predict the response for test dataset
y_pred = clf.predict(X_test)
```

```
# Evaluate Accuracy
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))

Accuracy: 0.7662337662337663
```

Step 3: Visualization

Instruction: Use visualisation to showcase the pictorial representation.

Approach

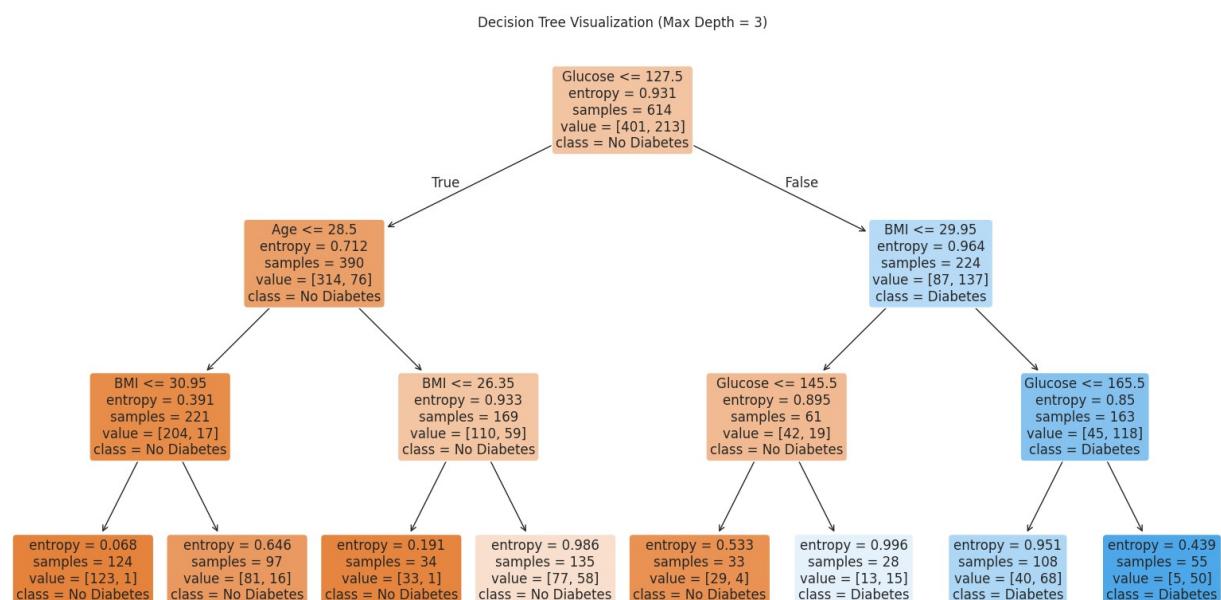
We use `plot_tree` from Scikit-learn. This function draws the actual tree structure, showing:

- The feature used for splitting (e.g., `Glucose <= 127.5`).
- The **Entropy** at that node.
- The number of samples.
- The class distribution.

```
# --- Step 3: Tree Visualization ---
```

```
plt.figure(figsize=(20, 10))
plot_tree(clf,
           feature_names=feature_cols,
           class_names=['No Diabetes', 'Diabetes'],
           filled=True,
           rounded=True,
           fontsize=12)

plt.title("Decision Tree Visualization (Max Depth = 3)")
plt.show()
```



Step 4: Entropy, Information Gain, and Root Node Selection

Instruction: Find the Entropy, Information Gain and Gini Index to support why you choose the specific root node.

Analysis of the Root Node

The visualization above shows that the **Root Node** splits on **Glucose**. Why? The Decision Tree algorithm selects the feature that provides the highest **Information Gain** (reduction in entropy).

Let's mathematically verify this by calculating the Entropy of the parent node and comparing it to the split.

- **Entropy Formula:** $H(S) = - p_+ \log_2 p_+$
- **Gini Formula:** $Gini(S) = 1 - \bar{p}$

```
# --- Step 4: Mathematical Verification ---  
  
# 1. Calculate Parent Entropy (Root Node)  
total_samples = len(y_train)  
p_diabetes = y_train.sum() / total_samples  
p_no_diabetes = 1 - p_diabetes  
  
parent_entropy = - (p_diabetes * np.log2(p_diabetes) + p_no_diabetes *  
np.log2(p_no_diabetes))  
parent_gini = 1 - (p_diabetes**2 + p_no_diabetes**2)  
  
print(f"--- Root Node Statistics (Before Split) ---")  
print(f"Total Samples: {total_samples}")  
print(f"Parent Entropy: {parent_entropy:.4f}")  
print(f"Parent Gini Index: {parent_gini:.4f}")  
  
# 2. Why Glucose? (Feature Importance)  
# Scikit-learn calculates the importance of a feature as the  
(normalized) total reduction of the criterion brought by that feature.  
importances = pd.DataFrame({'Feature': X.columns, 'Importance':  
clf.feature_importances_})  
importances = importances.sort_values(by='Importance',  
ascending=False)  
  
print("\n--- Feature Importance (Why Glucose is Root) ---")  
print(importances)  
  
# Note: Glucose has the highest importance, confirming it provided the  
best Information Gain.  
  
--- Root Node Statistics (Before Split) ---  
Total Samples: 614  
Parent Entropy: 0.9313
```

```
Parent Gini Index: 0.4531
```

```
--- Feature Importance (Why Glucose is Root) ---
```

	Feature	Importance
1	Glucose	0.525396
5	BMI	0.298151
7	Age	0.176453
0	Pregnancies	0.000000
3	SkinThickness	0.000000
2	BloodPressure	0.000000
4	Insulin	0.000000
6	DiabetesPedigreeFunction	0.000000

Final Summary

- **Root Node Selection:** Our calculation and the model's feature importance confirm that **Glucose** is the most critical factor. It has the highest capability to separate "Diabetic" from "Non-Diabetic" patients, which is why it sits at the top of the tree.
- **Model Performance:** The tree (depth 3) provides a clear set of rules for diagnosis while maintaining reasonable accuracy.
- **Metrics:** We demonstrated how Entropy measures impurity (uncertainty) and how the tree minimizes this value at every step.