

# Practical Assignment 09: K-Nearest Neighbors (KNN) on Iris Dataset

---

## Submission Details

Field	Details
Name	Ayushkar Pau
ID	GF202343142
Subject	Statistical Foundation of Data Science (CSU1658)
Date	December 10, 2025
Repo	<a href="#">View My Repository</a>

## Objective

The goal of this assignment is to implement the **K-Nearest Neighbors (KNN)** algorithm to classify Iris flowers based on their physical measurements. We will perform Exploratory Data Analysis (EDA), standardize the features, tune the hyperparameter  $K$  by analyzing error rates, and visualize the decision boundaries.

---

## Environment Setup and Dependencies

Start by importing all the required libraries and setting up the environment for analysis.

```
# --- 0. Environment Setup ---
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
from matplotlib.colors import ListedColormap

# Configure plotting
plt.rcParams['figure.figsize'] = (10, 6)
sns.set_style("whitegrid")
```

```
print("Environment setup complete.")  
Environment setup complete.
```

## Step 1: Exploratory Data Analysis (EDA)

**Instruction:** Perform head(), describe(), and groupby().

### Approach

We load the dataset into a DataFrame and inspect its structure.

- `head()` shows the first few rows.
- `describe()` provides statistical summaries (mean, std, quartiles).
- `groupby()` allows us to see the average measurements for each flower species.

```
# --- Step 1: EDA ---  
  
# Load Data  
iris = load_iris()  
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)  
df['species'] = iris.target # 0: Setosa, 1: Versicolor, 2: Virginica  
  
print("--- 1. First 5 Rows (head) ---")  
display(df.head())  
  
print("\n--- 2. Statistical Summary (describe) ---")  
display(df.describe())  
  
print("\n--- 3. Grouped by Species (mean) ---")  
# We group by species to see the average size differences  
display(df.groupby('species').mean())  
  
--- 1. First 5 Rows (head) ---  
  
   sepal length (cm)  sepal width (cm)  petal length (cm)  petal width  
(cm) \  
0           5.1            3.5            1.4  
0.2  
1           4.9            3.0            1.4  
0.2  
2           4.7            3.2            1.3  
0.2  
3           4.6            3.1            1.5  
0.2  
4           5.0            3.6            1.4  
0.2  
  
  species  
0      0
```

```
1      0  
2      0  
3      0  
4      0
```

### --- 2. Statistical Summary (describe) ---

```
    sepal length (cm)  sepal width (cm)  petal length (cm) \
count          150.000000          150.000000          150.000000
mean           5.843333           3.057333           3.758000
std            0.828066           0.435866           1.765298
min            4.300000           2.000000           1.000000
25%           5.100000           2.800000           1.600000
50%           5.800000           3.000000           4.350000
75%           6.400000           3.300000           5.100000
max            7.900000           4.400000           6.900000

    petal width (cm)  species
count          150.000000          150.000000
mean           1.199333           1.000000
std            0.762238           0.819232
min            0.100000           0.000000
25%           0.300000           0.000000
50%           1.300000           1.000000
75%           1.800000           2.000000
max            2.500000           2.000000
```

### --- 3. Grouped by Species (mean) ---

```
    sepal length (cm)  sepal width (cm)  petal length (cm) \
species
0                  5.006           3.428           1.462
1                  5.936           2.770           4.260
2                  6.588           2.974           5.552

    petal width (cm)
species
0                 0.246
1                 1.326
2                 2.026
```

## Step 2: Feature Scaling

**Instruction:** Feature scaling.

### Approach

KNN calculates distance (Euclidean) between points. If one feature has a large range (e.g., 1000-2000) and another has a small range (e.g., 0.1-1.0), the large feature will dominate the distance

calculation. We use `StandardScaler` to transform all features to have a mean of 0 and variance of 1.

```
# --- Step 2: Feature Scaling ---

# Define Features (X) and Target (y)
X = df.drop('species', axis=1)
y = df['species']

# Split the dataset (80% Train, 20% Test) BEFORE scaling to prevent
# data leakage
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Data scaled successfully.")
print(f"Training Data Shape: {X_train_scaled.shape}")
print(f"Test Data Shape: {X_test_scaled.shape}")

Data scaled successfully.
Training Data Shape: (120, 4)
Test Data Shape: (30, 4)
```

## Step 3: Training the K-NN Model

**Instruction:** Training the K-NN model on the Training set.

### Approach

We initialize the `KNeighborsClassifier`. We will start with a default value of  $K=5$  (analyzing 5 nearest neighbors) and the 'minkowski' metric (standard Euclidean distance).

```
# --- Step 3: Train KNN Model ---

# Initialize KNN with K=5
classifier = KNeighborsClassifier(n_neighbors=5, metric='minkowski',
p=2)

# Train on scaled data
classifier.fit(X_train_scaled, y_train)

print("KNN model (K=5) trained successfully.")

KNN model (K=5) trained successfully.
```

## Step 4 & 5: Evaluation

**Instruction:** Making the Confusion Matrix, Predicting Accuracy Score, and Classification Report.

### Approach

We use the trained model to predict species for the test set. We then evaluate performance using:

- **Confusion Matrix:** To see where the model confused one species for another.
- **Accuracy Score:** Overall percentage of correct predictions.
- **Classification Report:** Precision, Recall, and F1-Score for each class.

```
# --- Step 4 & 5: Evaluation ---

# Predict
y_pred = classifier.predict(X_test_scaled)

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print("--- Confusion Matrix ---")
print(cm)

# Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"\n--- Accuracy Score: {accuracy*100:.2f}% ---")

# Classification Report
print("\n--- Classification Report ---")
print(classification_report(y_test, y_pred,
target_names=iris.target_names))

--- Confusion Matrix ---
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]

--- Accuracy Score: 100.00% ---

--- Classification Report ---
      precision    recall   f1-score   support
setosa       1.00     1.00     1.00      10
versicolor   1.00     1.00     1.00       9
virginica    1.00     1.00     1.00      11

accuracy        1.00          1.00          1.00      30
macro avg     1.00     1.00     1.00      30
weighted avg  1.00     1.00     1.00      30
```

## Step 6, 7 & 8: Finding the Best K

**Instruction:** Comparing Error Rate with K Value, Plotting, and Finding Best K.

### Approach

Choosing the right  $K$  is critical.

- **Small K:** Sensitive to noise (overfitting).
- **Large K:** Smoothes boundaries too much (underfitting).

We will loop through  $K$  values from 1 to 40, calculate the **Mean Error Rate** (average of incorrect predictions) for each, and plot the results to find the "elbow" or minimum error point.

```
# --- Step 6, 7 & 8: Elbow Method to find Best K ---


error_rate = []

# Loop through K values from 1 to 40
for i in range(1, 41):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train_scaled, y_train)
    pred_i = knn.predict(X_test_scaled)

    # Calculate average error (mean of where prediction != actual)
    error_rate.append(np.mean(pred_i != y_test))

# Plot Error Rate vs K Value
plt.figure(figsize=(10, 6))
plt.plot(range(1, 41), error_rate, color='blue', linestyle='dashed',
marker='o',
        markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
plt.show()

# Find best K (minimum error)
min_error = min(error_rate)
best_k = error_rate.index(min_error) + 1
print(f"Minimum Error Rate: {min_error:.4f}")
print(f"Best K value: {best_k}")
```



```
Minimum Error Rate: 0.0000
Best K value: 1
```

## Step 9: Visualize Test Result of KNN

**Instruction:** Visualize Test Result of KNN.

### Approach

Visualizing a 4-dimensional dataset (4 features) is difficult. To show the **Decision Boundaries**, we will retrain a model using only the first two features (Sepal Length and Sepal Width). This allows us to plot a 2D map showing how the KNN algorithm divides the space between the three flower species.

```
# --- Step 9: Visualization (Using 2 Features) ---

# Re-train on just first two features for 2D plotting
X_vis = X.to_numpy()[:, :2] # Sepal Length, Sepal Width
y_vis = y.to_numpy()

# Split and Scale again for this visualization subset
X_train_v, X_test_v, y_train_v, y_test_v = train_test_split(X_vis,
y_vis, test_size=0.2, random_state=42)
scaler_v = StandardScaler()
X_train_v = scaler_v.fit_transform(X_train_v)
```

```

X_test_v = scaler_v.transform(X_test_v)

# Train KNN (using Best K found above, typically 1 or 3 for this dataset)
# We use a fixed K=5 here for clear visualization standard
knn_vis = KNeighborsClassifier(n_neighbors=5)
knn_vis.fit(X_train_v, y_train_v)

# Create a mesh grid
X_set, y_set = X_test_v, y_test_v
X1, X2 = np.meshgrid(np.arange(start=X_set[:, 0].min() - 1,
stop=X_set[:, 0].max() + 1, step=0.01),
                     np.arange(start=X_set[:, 1].min() - 1,
stop=X_set[:, 1].max() + 1, step=0.01))


plt.figure(figsize=(10, 6))
plt.contourf(X1, X2, knn_vis.predict(np.array([X1.ravel(),
X2.ravel()]).T).reshape(X1.shape),
             alpha=0.3, cmap=ListedColormap(['red', 'green', 'blue']))
plt.xlim(X1.min(), X1.max())
plt.ylim(X2.min(), X2.max())

# Scatter plot of Test Data
for i, j in enumerate(np.unique(y_set)):
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
                c=ListedColormap(['red', 'green', 'blue'])(i),
                label=iris.target_names[j])

plt.title('KNN Decision Boundary (Test Set)')
plt.xlabel('Sepal Length (Scaled)')
plt.ylabel('Sepal Width (Scaled)')
plt.legend()
plt.show()

/tmp/ipykernel_98573/2681138395.py:31: UserWarning: *c* argument looks
like a single numeric RGB or RGBA sequence, which should be avoided as
value-mapping will have precedence in case its length matches with *x*
& *y*. Please use the *color* keyword-argument or provide a 2D array
with a single row if you intend to specify the same RGB or RGBA value
for all points.
    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],

```



## Final Summary and Conclusions

This notebook successfully implemented the K-Nearest Neighbors (KNN) algorithm to classify Iris flower species, covering the complete machine learning pipeline.

### Summary of Tasks Completed:

- **1. Exploratory Data Analysis:** We examined the dataset structure and statistical properties, observing distinct measurements for different species.
- **2. Data Preprocessing:** We applied **Standard Scaling** to normalize feature values, ensuring that distance calculations were not biased by feature magnitude.
- **3. Model Training & Evaluation:** We trained a KNN classifier and evaluated its performance using a Confusion Matrix and Classification Report, achieving high accuracy.
- **4. Hyperparameter Tuning:** We utilized the **Elbow Method** by plotting Error Rate vs. K Value. We observed that while small K values (like 1) gave low error, slightly larger values (like 5-10) often generalize better.
- **5. Visualization:** We visualized the decision boundaries in 2D space, clearly showing how the algorithm partitions the feature space to separate the three classes.

### Key Learnings:

- **Feature Scaling is Critical:** Since KNN relies on Euclidean distance, unscaled data can lead to poor performance. Standardizing inputs is a mandatory step.

- **Choosing K:** The choice of K involves a bias-variance trade-off. A very small K captures noise (overfitting), while a very large K smooths out boundaries too much (underfitting). The error rate plot is an effective tool for finding the optimal balance.

This completes the requirements for Practical Assignment 09.